

# Deep Neural Networks

Project week9:  
Adding BatchNormalization

# Initial model: 6 classes, not too complex

```
model.summary()
```

Model: "sequential\_13"

Layer (type)	Output Shape	Param #
zero_padding2d_11 (ZeroPaddi	(None, 226, 226, 3)	0
conv2d_30 (Conv2D)	(None, 224, 224, 16)	448
max_pooling2d_30 (MaxPooling	(None, 112, 112, 16)	0
conv2d_31 (Conv2D)	(None, 110, 110, 32)	4640
max_pooling2d_31 (MaxPooling	(None, 109, 109, 32)	0
flatten_13 (Flatten)	(None, 380192)	0
dropout_13 (Dropout)	(None, 380192)	0
dense_13 (Dense)	(None, 6)	2281158
Total params: 2,286,246		
Trainable params: 2,286,246		
Non-trainable params: 0		

Training accuracy: 0.83  
Validation accuracy: 0.88  
Test accuracy: 0.83

# More layers

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
zero_padding2d_14 (ZeroPaddi	(None, 226, 226, 3)	0
conv2d_38 (Conv2D)	(None, 224, 224, 16)	448
max_pooling2d_38 (MaxPooling	(None, 112, 112, 16)	0
conv2d_39 (Conv2D)	(None, 110, 110, 32)	4640
max_pooling2d_39 (MaxPooling	(None, 109, 109, 32)	0
conv2d_40 (Conv2D)	(None, 107, 107, 16)	4624
max_pooling2d_40 (MaxPooling	(None, 53, 53, 16)	0
flatten_16 (Flatten)	(None, 44944)	0
dropout_16 (Dropout)	(None, 44944)	0
dense_16 (Dense)	(None, 6)	269670
Total params: 279,382		
Trainable params: 279,382		
Non-trainable params: 0		

[+ Code](#)[+ Markdown](#)

Training accuracy: 0.97  
Validation accuracy: 0.97  
Test accuracy: 0.98

# Adding BatchNormalization

Model: "sequential\_18"

Layer (type)	Output Shape	Param #
zero_padding2d_16 (ZeroPaddi	(None, 226, 226, 3)	0
conv2d_44 (Conv2D)	(None, 224, 224, 16)	448
batch_normalization_17 (Batc	(None, 224, 224, 16)	64
max_pooling2d_44 (MaxPooling	(None, 112, 112, 16)	0
conv2d_45 (Conv2D)	(None, 110, 110, 32)	4640
batch_normalization_18 (Batc	(None, 110, 110, 32)	128
max_pooling2d_45 (MaxPooling	(None, 109, 109, 32)	0
conv2d_46 (Conv2D)	(None, 107, 107, 16)	4624
batch_normalization_19 (Batc	(None, 107, 107, 16)	64
max_pooling2d_46 (MaxPooling	(None, 53, 53, 16)	0
flatten_18 (Flatten)	(None, 44944)	0
dropout_18 (Dropout)	(None, 44944)	0
dense_18 (Dense)	(None, 6)	269670
Total params: 279,638		
Trainable params: 279,510		
Non-trainable params: 128		

Training accuracy: 0.99  
Validation accuracy: 0.44  
Test accuracy: 0.46

# Why is batch normalization reducing my model training accuracy?

Asked 1 year, 8 months ago   Modified 1 year, 8 months ago   Viewed 702 times

1 Answer

Sorted by: Highest score (default)



2



Batch Normalisation doesn't guarantee that your performance will increase. But it does work well in some cases.

One of things you can try to do is:

1. Increase the batch size of the training. This will give a more appropriate mean and standard deviation for normalisation.
2. Play around with the BN parameters, specifically the momentum parameter. See more here about the params [https://keras.io/api/layers/normalization\\_layers/batch\\_normalization/](https://keras.io/api/layers/normalization_layers/batch_normalization/) I would suggest to decrease the momentum and try again.
3. If it still doesn't work, leave it out.

[Share](#) [Follow](#)

answered Mar 4, 2021 at 8:26



s510

1,741 ● 6 ● 17

# Parameters in BatchNormalization

## BatchNormalization layer

**BatchNormalization** class

```
tf.keras.layers.BatchNormalization(  
    axis=-1,  
    momentum=0.99,  
    epsilon=0.001,  
    center=True,  
    scale=True,  
    beta_initializer="zeros",  
    gamma_initializer="ones",  
    moving_mean_initializer="zeros",  
    moving_variance_initializer="ones",  
    beta_regularizer=None,  
    gamma_regularizer=None,  
    beta_constraint=None,  
    gamma_constraint=None,  
    **kwargs  
)
```

- [https://keras.io/api/layers/normalization\\_layers/batch\\_normalization/](https://keras.io/api/layers/normalization_layers/batch_normalization/)

# Recap: why BatchNormalization?

- Issues without BatchNormalization
  - Internal covariate shift
    - update of weights → slight changes → input distribution of every neuron is different at every epoch
    - in deep networks: add up fast, amplify greatly
    - → these neurons need to continuously adapt to the changing input distribution, meaning that their learning capabilities are severely bottlenecked
  - Vanishing and exploding gradients when using larger learning rates
- Advantages:
  - Accelerate deep network training by reducing internal covariate shift
  - can use larger learning rates
  - Reduces overfitting
    - regularising effect since it adds noise to the inputs of every layer

# BatchNormalization during training vs. test phase

- During training:
  - Mean and standard deviation calculated using samples in the mini-batch
- During testing:
  - Does not make sense to calculate new values
  - → use a running mean and running variance that is calculated during training
  - Parameter: momentum



# “momentum”

```
running_mean = momentum * running_mean + (1-momentum) * new_mean  
running_var = momentum * running_var + (1-momentum) * new_var
```

Momentum is the importance given to the last seen mini-batch, a.k.a “lag”. If the momentum is set to 0, the running mean and variance come from the last seen mini-batch. However, this may be biased and not the desirable one for testing. Conversely, if momentum is set to 1, it uses the running mean and variance from the first mini-batch. Essentially, momentum controls how much each new mini-batch contributes to the running averages.

Ideally, the momentum should be set close to 1 ( $>0.9$ ) to ensure slow learning of the running mean and variance such that the noise in a mini-batch is ignored.

# Final model

```
[564]: # define the keras model
model = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.ZeroPadding2D(),
        layers.Conv2D(16, kernel_size=(3, 3), activation="relu"),
        layers.BatchNormalization(momentum=0.8),
        layers.MaxPooling2D(pool_size=(2, 2), strides=2),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.BatchNormalization(momentum=0.8),
        layers.MaxPooling2D(pool_size=(2, 2), strides=1),
        layers.Conv2D(16, kernel_size=(3, 3), activation="relu"),
        layers.BatchNormalization(momentum=0.8),
        layers.MaxPooling2D(pool_size=(2, 2), strides=2),
        layers.Flatten(),
        #layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)
```

Training accuracy: 1  
Validation accuracy: 0.99  
Test accuracy: 0.99

- <https://towardsdatascience.com/batch-normalisation-explained-5f4bd9de5feb?gi=21f16071fb92>